

Basics

Arrays

Contiguous storage: elements stored sequentially in memory.

$O(1)$ access: direct access by index.

Memory Layout:

Index:	0	1	2	3
Value:	2	7	11	15
	↑	↑	↑	↑
	arr[0]	arr[1]	arr[2]	arr[3]

Drawback: insertion/deletion requires shifting elements → **$O(n)$** .

Before:

Index:	0	1	2	3
Value:	2	7	11	15

Insert 99 at index 2 → shift [11, 15] to the right

After:

Index:	0	1	2	3	4
Value:	2	7	99	11	15

Before:

Index:	0	1	2	3
Value:	2	7	11	15

Delete 7 at index 1 → shift [11, 15] to the left

After:

Index:	0	1	2
Value:	2	11	15

Use case:

- efficient for sequential traversal;
- inefficient for frequent middle insertions/deletions.

Hash Tables

Now let's talk about hash tables.

A **hash table** is a general data structure.

The core idea is to use a hash function to map a key into an array index, so you can quickly store or retrieve the value.

Hash Function: $\text{key} \rightarrow \text{index}$

Keys:	[apple]	[banana]	[cat]
Index:	2	5	8
Value:	"red"	"yellow"	"black"

Advantage: $O(1)$ average time for lookup, insertion, deletion

- **HashMap:** stores $\text{key} \rightarrow \text{value}$ (Python: `dict`)
- **HashSet:** stores only keys (Python: `set`), used for deduplication / membership check

In algorithm problems, hash tables are most commonly used for two things:

1. **Counting frequency** – find most frequent elements, detect duplicates

```
map[num] += 1
```

2. **Recording positions** – quick lookup of element's index

```
map[num] = index
```

So overall, the hash table gives us:

- Quick existence checks
- Quick frequency counts
- Quick position lookups

Arrays vs. Hash Tables

So here's the key question: when is an array alone enough, and when should we bring in a hash table?

Here's a simple way to decide:

If the problem is about **sequential processing**, like maximum subarray sum or sliding window maximum, an array by itself is usually enough.

But if the problem mentions things like:

- “check quickly if something has appeared before,”
- “count frequencies,”
- “find the position of a target value,”
- or explicitly “achieve $O(1)$ time complexity,”

then it's time to consider a hash table.

Example: Two Sum

Given an array of integers `nums` and an integer `target`, return the **indices** of the two numbers such that they add up to `target`.

You may assume that each input has exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Example:

```
Input: nums = [2,7,11,15], target = 9
```

```
Output: [0,1]
```

```
Explanation: Because nums[0] + nums[1] == 9, we return [0, 1].
```

Brute Force

Idea: try every possible pair

1. Outer loop picks the first number
2. Inner loop picks the second number
3. If their sum equals `target`, return indices

Time complexity: $O(n^2)$ → inefficient for large arrays

```
def twoSum(nums, target):
    n = len(nums)
    for i in range(n):
        for j in range(i + 1, n): # start j from i+1 to avoid duplicates
            if nums[i] + nums[j] == target:
                return [i, j]
```

Optimized with a Hash Table

Store mapping **value** → **index** while traversing

1. For each number `num`, compute `complement = target - num`
2. Check if `complement` exists in hash table
3. If yes → return `[hashmap[complement], i]`
4. If no → store current `num: i` in hash table

Each number processed once, each lookup $O(1)$. Total time complexity: **$O(n)$** .

```
def twoSum(nums, target):
    hashmap = {} # key = value, value = index
    for i, num in enumerate(nums):
        complement = target - num
        if complement in hashmap:
            return [hashmap[complement], i]
        hashmap[num] = i
```

Why Store Indices?

Now, here's an important detail:

If the question only asked, *“Does the array contain two numbers that add up to target?”* Then we don't need indices at all. We could just use a **hash set** to check if `target - num` has appeared before, and return True if it has.

But this problem specifically asks for **indices**. That's why we need a **hash map**, not just a hash set. The hash map doesn't just store values — it stores their positions. This way, when we find a complement, we can return the exact pair of indices.

Handy Libraries: Counter & defaultdict

In Python, a hash table is a `dict`, but using plain `dict` can be verbose:

- Counting frequencies → need to check key existence before updating.
- Grouping items → need to initialize empty lists manually.

```
# Using a normal dict to count frequencies
nums = [1, 1, 2, 3, 3, 3]
count = {}

for num in nums:
    if num not in count:    # need to check manually
        count[num] = 0
    count[num] += 1

print(count)    # {1: 2, 2: 1, 3: 3}
```

To simplify:

- `Counter`: one-line frequency counting.
- `defaultdict`: automatic initialization of default values (e.g., `int`, `list`, `set`).

These tools reduce boilerplate and make code cleaner.

`collections.Counter`

`Counter` lets you count frequencies in one line, instead of writing loops and conditionals yourself.

```
from collections import Counter

nums = [1,1,2,3,3,3]
counter = Counter(nums)
print(counter)    # Output: Counter({3: 3, 1: 2, 2: 1})
```

Things to keep in mind:

- Subclass of `dict` → implemented as a hash table, **O(n)** time.
- Usable like a normal dict (`counter[key]`). The only difference is the print format looks like `Counter({...})`.
- Missing keys default to **0** (no `KeyError`).

- Supports arithmetic & set-like operations:
 - `Counter(a) + Counter(b)` → add counts
 - `Counter(a) - Counter(b)` → subtract counts (removes ≤ 0)
- Slight overhead due to extra features → for large datasets, `dict` or `defaultdict(int)` may be faster.

collections.defaultdict

Normally with a `dict`, if a key doesn't exist, you have to do:

```
if key not in dict:
    dict[key] = []
```

With `defaultdict`, you don't need this check. It automatically initializes a default value for you.

```
from collections import defaultdict

groups = defaultdict(list)
words = ["eat", "tea", "tan", "ate", "nat", "bat"]

for w in words:
    key = "".join(sorted(w))
    groups[key].append(w)

print(groups)
# Output: {'aet': ['eat', 'tea', 'ate'], 'ant': ['tan', 'nat'], 'abt': ['bat']}
```

Things to keep in mind:

- You must pass a **factory function** when creating it:
 - `defaultdict(int)` → default is 0
 - `defaultdict(list)` → default is []
 - `defaultdict(set)` → default is set()
 - You **cannot** pass a raw value like `defaultdict(0)` — that will raise an error.
- It **automatically creates keys** when accessed:

```
from collections import defaultdict
d = defaultdict(int)
print(d[1])      # prints 0, and also creates key=1
print(d)         # defaultdict(<class 'int'>, {1: 0})
```

- Printing looks like `defaultdict(<class 'list'>, {...})`, not a plain dict. If you need to return the result to a platform like LeetCode, better convert it: `dict(d)`.
- Performance-wise, it's the same as a normal dict — $O(1)$ for insert/lookup.

Common Problem Types

Today's Problems

- [1. Two Sum](#)
- [128. Longest Consecutive Sequence](#)
- [36. Valid Sudoku](#)

Hash Table Counting + Array Traversal

Core idea: first use a hash table to count frequencies or features, then traverse the array to do grouping or pairing.

- [Group Anagrams](#)
- [Top K Frequent Elements](#)

Two Pointers + Hash Table

Core idea: use two pointers to maintain a sliding window, and use a hash table to keep track of element frequencies inside the window. This way you can dynamically expand/shrink the window while quickly checking conditions.

- [3. Longest Substring Without Repeating Characters](#)

Subarray / Substring Problems

Core idea: use a hash table to store element positions or prefix-sum counts, then traverse the array.

- [560. Subarray Sum Equals K](#)

- [523. Continuous Subarray Sum](#)

Design Problems

These test whether you can use a hash table to implement your own data structure.

- [380. Insert Delete GetRandom O\(1\)](#)